

# Equalizer: A Mature Parallel Rendering Framework

Stefan Eilemann\*

Renato Pajarola†

Visualization and MultiMedia Lab  
 Department of Informatics  
 University of Zürich

**Abstract**— We present the features, algorithms and system integration necessary to implement a parallel rendering framework usable in a wide range of real-world research and industry applications, based on the basic architecture and implementation of the Equalizer parallel rendering framework presented in [4].

**Index Terms**—Parallel Rendering, Scalable Visualization, Cluster Graphics, Immersive Environments, Display Walls

## 1 INTRODUCTION

The continuing improvements in hardware integration lead to ever faster CPUs and GPUs, as well as higher resolution sensor and display devices. Moreover, increased hardware parallelism is applied in form of multi-core CPU workstations, massive parallel super computers, or cluster systems. Hand in hand goes the rapid growth in complexity of data sets from numerical simulations, high-resolution 3D scanning systems or bio-medical imaging, which causes interactive exploration and visualization of such large data sets to become a serious challenge. It is thus crucial for a visualization solution to take advantage of hardware accelerated scalable parallel rendering. In this systems paper we describe a new scalable parallel rendering framework called *Equalizer* that is aimed primarily at cluster-parallel rendering, but works as well in a shared-memory system. Cluster systems are the main focus because workstation graphics hardware is developing faster than high-end graphics (super-) computers can absorb new developments, and also because clusters offer a better cost-performance balance.

Previous parallel rendering approaches typically failed in one of the following system requirements:

- a) generic application support, instead of special domain solution
- b) scalable abstraction of the graphics layer
- c) exploit existing code infrastructure, such as proprietary scene graphs, molecular data structures, level-of-detail and geometry databases

To date, generic and scalable parallel rendering frameworks that can be adopted to a wide range of scientific visualization domains are not yet readily available. Furthermore, flexible configurability to arbitrary cluster and display-wall configurations has also not been addressed in the past, but is of immense practical importance to scientists depending high-performance interactive visualization as a scientific tool. In this paper we present Equalizer, which is a novel flexible framework for parallel rendering that supports scalable performance, configuration flexibility, is *minimally invasive* with respect to adapting existing visualization applications, and is applicable to virtually any scientific visualization application domain.

The main contributions that Equalizer introduces in a single parallel rendering system, and which are presented in this paper are:

- i) novel concept of compound trees for flexible configuration of graphics system resources,
- ii) easy specification of parallel task decomposition and image compositing choice through compound tree layouts,

\*email: eilemann@gmail.com

†email: pajarola@acm.org

*Manuscript received 5 February 2008; accepted ? April 2008; posted online ? June 2008.*

*For information on obtaining reprints of this article, please send e-mail to:  
 tvcg@computer.org.*

- iii) automatic decomposition and distributed execution of rendering tasks according to compound tree,
- iv) support for parallel surface as well as transparent (volume) rendering through z-visibility as well as  $\alpha$ -blending compositing,
- v) fully decentralized architecture providing network swap barrier (synchronization) and distributed objects functionality,
- vi) support for low-latency distributed frame synchronization and image compositing,
- vii) minimally invasive programming model.

Equalizer is open source, available under the LGPL license from <http://www.equalizergraphics.com/>, which allows it to be used both for open source and commercial applications. It is source-code portable, and has been tested on Linux, Microsoft Windows, and Mac OS X in 32 and 64 bit mode using both little endian and big endian processors.

## 2 RELATED WORK

[4] summarized the work in parallel rendering at the time of publication. In the following we will present the related work published since this publication.

[5] cross-segment load balancing

[2]: Framework to develop apps for multitile, thus targeted specifically for tiled-wall visualization systems: grid configuration, scalability (?), multiple work sessions, multiuser events. Master-slave approach like Eq. with grid nodes running instances of the application. Is maybe more like a GLUT and window replacement, or enhancement to deal with the OpenGL contexts and events

[1] tiled display wall virtual environment

[?] DisplayCluster

[7] ClusterGL?

Omegalib

## 3 USABILITY

In this section we present features motivated by real-world application use cases, i.e., new functionalities rather than performance improvements. We motivate the use case, explain the architecture and integration into our parallel rendering framework, and, where applicable, show the steps needed to use this functionality in applications.

### 3.1 Physical and Logical Visualization Setup

Real-world visualization setups are often complex, and having an abstract representation of the display system can simplify the configuration process. Real-world applications often have the need to be aware of spatial relationship of the display setup, for example to render 2D overlays or to configure multiple views on a tiled display wall.

We addressed this need through a new configuration section interspersed between the node/pipeline/window/channel hardware resources and the compound trees configurating the resource usage for parallel rendering.

A typical installation consists of one projection canvas, which is one aggregated projection surface, e.g., a tiled display wall or a CAVE. Desktop windows are considered a canvas. Each canvas is made of one or more segments, which are the individual outputs connected to

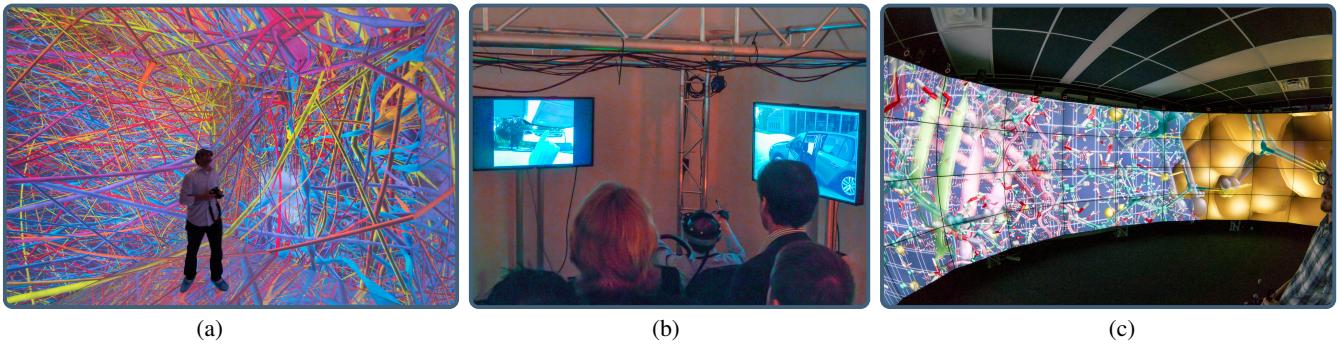


Fig. 1. Example Equalizer application use cases: (a) 192 Megapixel CAVE at KAUST running RTNeuron, (b) Immersive HMD with external tracked and untracked views running RTT DeltaGen for virtual car usability studies (c) Cave2 running a molecular visualization build using Omegalib.

a display or projector. Segments can be planar or non-planar to each other, and can overlap or have gaps between each other. A segment is referencing a channel, which defines the output area of this segment, e.g., on a DVI connector connected to a projector.

A canvas can define a frustum, which will create default, planar sub-frusta for all of its segments. A segment can also define a frustum, which overrides the canvas frustum, e.g., for non-planar setups such as CAVEs or curved screens. These frusta describe a physically-correct display setup for a Virtual Reality installation. A canvas may have a software or hardware swap barrier, which will synchronize the rendering of all contributing GPUs.

On each canvas, the application can display one or more views. A view is a view on a model, in the sense used by the MVC pattern. The view class is used by Equalizer applications to define view-specific data for rendering, e.g., a scene, viewing mode or camera. The application process manages this data, and the render clients receive it for rendering.

A layout groups one or more views which logically belong together. A layout is applied to a canvas. The layout assignment can be changed at run-time by the application. The intersection between views and segments defines which output channels are available, and which frustum they should use for rendering. These output channels are then used as destination channels in a compound. They are automatically created during configuration.

A view may have a frustum description. The view's frustum overrides frusta specified at the canvas or segment level. This is used for non-physically correct rendering, e.g., to compare two models side-by-side on a tiled display wall. If the view does not specify a frustum, the corresponding destination channels will use the physically correct sub-frustum resulting from the view/segment intersection.

An observer looks at one or more views. It is described by the observer position in the world and its eye separation. Each observer will have its own stereo mode, focus distance and frame loop (framerate). This allows to have untracked views and multiple tracked views, e.g., two HMDs, in the same application.

### 3.2 Automatic Configuration

Automatic configuration implements the discovery of local and remote resources as well as the creation of typical configurations using the discovered resources at application launch time.

The discovery is implemented in a separate library, hwsd (Hardware Service Discovery), which uses a plugin-based approach to discover GPUs for GLX, AGL or WGL windowing systems, as well as network interfaces on Linux, Mac OS X and Windows. Furthermore, it detects the presence of VirtualGL to allow optimal configuration of remote visualization clusters. The resources can be discovered on the local workstation, and through the help of a simple daemon using the zeroconf protocol, on a set of remote nodes within a visualization cluster. A session identifier may be used to support multiple users on a single cluster.

The Equalizer server uses the hwsd library to discover local and remote resources when an hwsd session name instead of a .eqc con-

figuration file is provided. A set of standard decomposition modes is configured, which can be selected through activating the corresponding layout.

This versatile mechanism allows non-experts to configure and profit from multi-GPU workstations and visualization clusters, as well as to provide system administrators with the tools to implement easy to use integration with cluster schedulers. This feature is transparent to Equalizer application developers.

### 3.3 Qt Windowing

Challenges in threading model, architecture

### 3.4 Tide Integration

Tiled interactive display environment, parallel pixel streaming, events

### 3.5 Sequel

application, renderer, view data

## 4 THE COLLAGE NETWORK LIBRARY

### 4.1 Distributed, Versioned Objects

types (instance, delta, static), versioning, multicast, compression, serializable with dirty bits, mapping, blocking commits

### 4.2 Reliable Stream Protocol

UDP-based reliability protocol

### 4.3 Infiniband RDMA

reverse-engr impl

## 5 VIRTUAL REALITY

### 5.1 Dynamic Focus Distance

Focus what user is looking at

### 5.2 Asymmetric Eye Position

Better HMD by measuring user geometry

### 5.3 Application-specific Scaling

Gullivers world

### 5.4 Runtime Stereo Switch

### 5.5 Swap Synchronization and GPU affinity

## 6 PERFORMANCE

### 6.1 New Decomposition Modes

The initial version of Equalizer implemented sort-first (2D), sort-last (DB) and stereo (EYE) decomposition. In the following we present new decomposition modes and motivate their use case.

#### 6.1.1 Time-Multiplex

Interaction with threading

### 6.1.2 Tiles and Chunks

Equalizers do what?

### 6.1.3 Pixel

Equal load for fill-limited apps

### 6.1.4 Subpixel

MSAA and DOF, plus idle refinement algo

## 6.2 Equalizers

Equalizer are an addition to compound trees. They modify parameters of their respective subtree at runtime to optimize one aspect of the decomposition.

### 6.2.1 Sort-First and Sort-Last Load Equalizer

reactive load equalization, describe tunables (damping, resistance, delta, granularity), sort-first respects input channel size

### 6.2.2 Cross-Segment Load Equalizer

summarize [5]

### 6.2.3 Dynamic Frame Resolution

Constant frame rate for fill-limited applications

### 6.2.4 Frame Rate Equalizer

Predominantly used by DPlex to smooth output framerate

### 6.2.5 Monitoring

Control workstation in VR setups

## 6.3 Optimizations

### 6.3.1 Region of Interest

The region of interest is the screen-space 2D bounding box enclosing the geometry rendered by a single resource. We have extended the core parallel rendering framework to use an application-provided ROI to optimize the load equalizer as well as image compositing performance. The load equalizer uses the ROI to refine its load grid to the regions containing data. The compositing code uses the ROI to minimize image readback and network transmission. In [3], we provide the details of the algorithm, and show that this improves rendering performance of up to 25%.

### 6.3.2 Asynchronous Compositing

Asynchronous compositing pipelines rendering with compositing operations, by executing the image readback, network transfer and image assembly from threads running in parallel to the rendering threads. In [3], we provide the details of the implementation and experimental data showing an improvement of the rendering performance of over 25% for large node counts.

### 6.3.3 Download and Compression Plugins

GPU-CPU transfer plugins with optional compression (eg YUV) linked to CPU compression for network transfer

### 6.3.4 Thread Synchronization Modes

Per-node sync, draw sync, async

## 7.4 Raster

## 7.5 Omegalib

## 8 EXPERIMENTAL RESULTS

### 8.1 Object Distribution

Distribute large ply model to 1..16 nodes using TCP, IB and RSP

### 8.2 Decomposition Modes

Two graphs: Time to render a 4k, 256-step MSAA image of largest ply model and Livre using 1..16 nodes with all modes (2D, 2D LB, DB, DPlex (over AA steps), tiles, chunks, pixel, subpixel)

## 9 DISCUSSION AND CONCLUSION

## ACKNOWLEDGEMENTS

We would like to thank and acknowledge the following institutions and projects for providing the 3D geometry and volume test data sets: the Digital Michelangelo Project, Stanford 3D Scanning Repository, Cyberware Inc., volvis.org and the Visual Human Project. This work was partially supported by the Swiss National Science Foundation Grant 200021-116329/1.

## REFERENCES

- [1] Y. Cho, M. Kim, and K. S. Park. LOTUS: Composing a multi-user interactive tiled display virtual environment. *The Visual Computer*, 28(1):99–109, 2012.
- [2] K.-U. Doerr and F. Kuester. CGLX: A scalable, high-performance visualization framework for networked display environments. *IEEE Transactions on Visualization and Computer Graphics*, 17(2):320–332, March 2011.
- [3] S. Eilemann, A. Bilgili, M. Abdellah, J. Hernando, M. Makhinya, R. Pajarola, and F. Schürmann. Parallel Rendering on Hybrid Multi-GPU Clusters. In *Proceedings of the 12th Eurographics Symposium on Parallel Graphics and Visualization*, pages 109–117, 2012.
- [4] S. Eilemann, M. Makhinya, and R. Pajarola. Equalizer: A scalable parallel rendering framework. *IEEE Transactions on Visualization and Computer Graphics*, 15(3):436–452, May/June 2009.
- [5] F. Erol, S. Eilemann, and R. Pajarola. Cross-segment load balancing in parallel rendering. In *Proceedings Eurographics Symposium on Parallel Graphics and Visualization*, pages 41–50, 2011.
- [6] J. B. Hernando, J. Biddiscombe, B. Bohara, S. Eilemann, and F. Schürmann. Practical Parallel Rendering of Detailed Neuron Simulations. In *Proceedings of the 13th Eurographics Symposium on Parallel Graphics and Visualization*, EGPGV, pages 49–56, Aire-la-Ville, Switzerland, Switzerland, 2013. Eurographics Association.
- [7] B. Neal, P. Hunkin, and A. McGregor. Distributed OpenGL rendering in network bandwidth constrained environments. In T. Kuhlen, R. Pajarola, and K. Zhou, editors, *Proceedings Eurographics Conference on Parallel Graphics and Visualization*, pages 21–29. Eurographics Association, 2011.